

# Variational AutoEncoder & Generative Models

By: Shai Harel, structured data vision team

# What we'll see today

- Generative vs. Discriminative models [1]
- VAE Algorithm Overview [2]
- Putting it to work - Semi-supervised [3]

[1] Deep Neural Networks are Easily Fooled

[2] Auto-Encoding Variational Bayes

[3] Semi-Supervised Learning with Deep Generative Models

# What we'll \*NOT\* see today

$$\log p(\mathbf{X}) = \sum_{i=1}^N \log p(\mathbf{x}^{(i)}), \quad (1)$$

but in general this marginal likelihood is intractable to compute or differentiate directly for flexible generative models that have high-dimensional latent variables and flexible priors and likelihoods. A solution is to introduce  $q(\mathbf{z}|\mathbf{x})$ , a parametric *inference model* defined over the latent variables, and optimize the *variational lower bound* on the marginal log-likelihood of each observation  $\mathbf{x}$ :

$$\log p(\mathbf{x}) \geq \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}, \mathbf{z}) - \log q(\mathbf{z}|\mathbf{x})] = \mathcal{L}(\mathbf{x}; \theta) \quad (2)$$

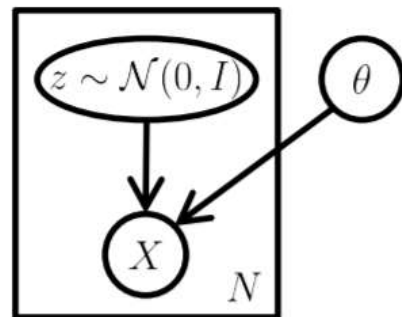
where  $\theta$  indicates the parameters of  $p$  and  $q$  models.

There are various ways to optimize the lower bound  $\mathcal{L}(\mathbf{x}; \theta)$ ; for continuous  $\mathbf{z}$  it can be done efficiently through a re-parameterization of  $q(\mathbf{z}|\mathbf{x})$  (Kingma & Welling, 2013; Rezende et al., 2014).

This way of optimizing the variational lower bound with a parametric inference network and re-parameterization of continuous latent variables is usually called VAE. The “autoencoding” terminology comes from the fact that the lower bound  $\mathcal{L}(\mathbf{x}; \theta)$  can be re-arranged:

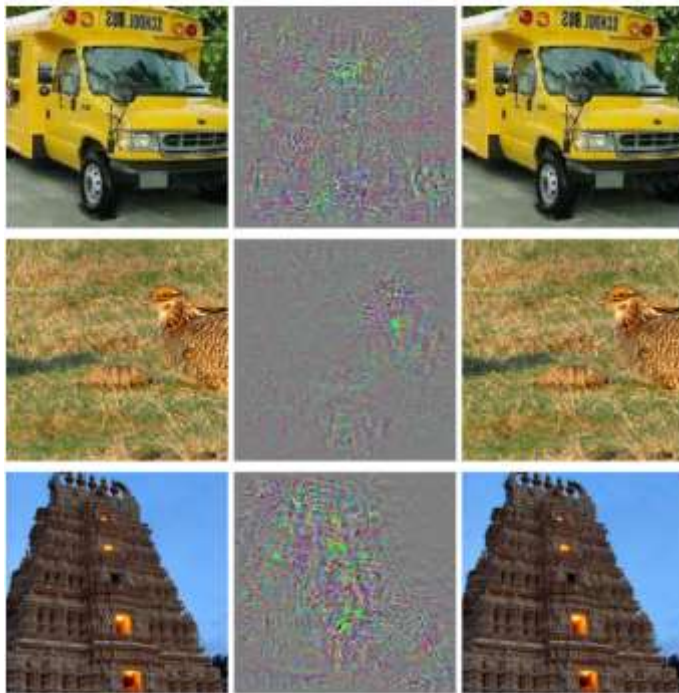
$$\mathcal{L}(\mathbf{x}; \theta) = \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}, \mathbf{z}) - \log q(\mathbf{z}|\mathbf{x})] \quad (3)$$

$$= \mathbb{E}_{q(\mathbf{z}|\mathbf{x})} [\log p(\mathbf{x}|\mathbf{z})] - D_{KL}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z})) \quad (4)$$



$$\begin{aligned} L &= \log(p(\mathbf{x})) \\ &= \sum_{\mathbf{z}} q(\mathbf{z}|\mathbf{x}) \log(p(\mathbf{x})) \\ &= \sum_{\mathbf{z}} q(\mathbf{z}|\mathbf{x}) \log\left(\frac{p(\mathbf{z}, \mathbf{x})}{p(\mathbf{z}|\mathbf{x})}\right) \\ &= \sum_{\mathbf{z}} q(\mathbf{z}|\mathbf{x}) \log\left(\frac{p(\mathbf{z}, \mathbf{x}) q(\mathbf{z}|\mathbf{x})}{q(\mathbf{z}|\mathbf{x}) p(\mathbf{z}|\mathbf{x})}\right) \\ &= \sum_{\mathbf{z}} q(\mathbf{z}|\mathbf{x}) \log\left(\frac{p(\mathbf{z}, \mathbf{x})}{q(\mathbf{z}|\mathbf{x})}\right) + \sum_{\mathbf{z}} q(\mathbf{z}|\mathbf{x}) \log\left(\frac{q(\mathbf{z}|\mathbf{x})}{p(\mathbf{z}|\mathbf{x})}\right) \\ &= L^v + D_{KL}(q(\mathbf{z}|\mathbf{x})||p(\mathbf{z}|\mathbf{x})) \\ &\geq L^v \end{aligned}$$

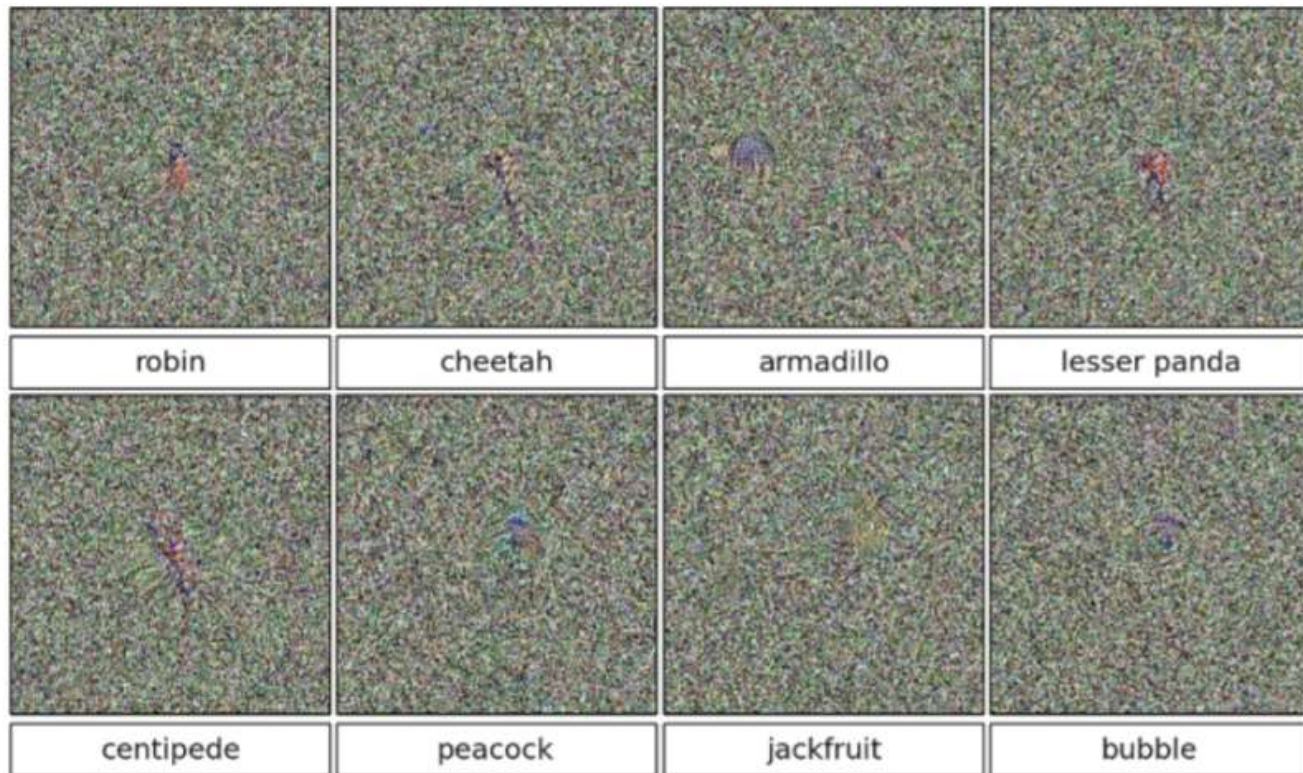
# Why Gen. Models?



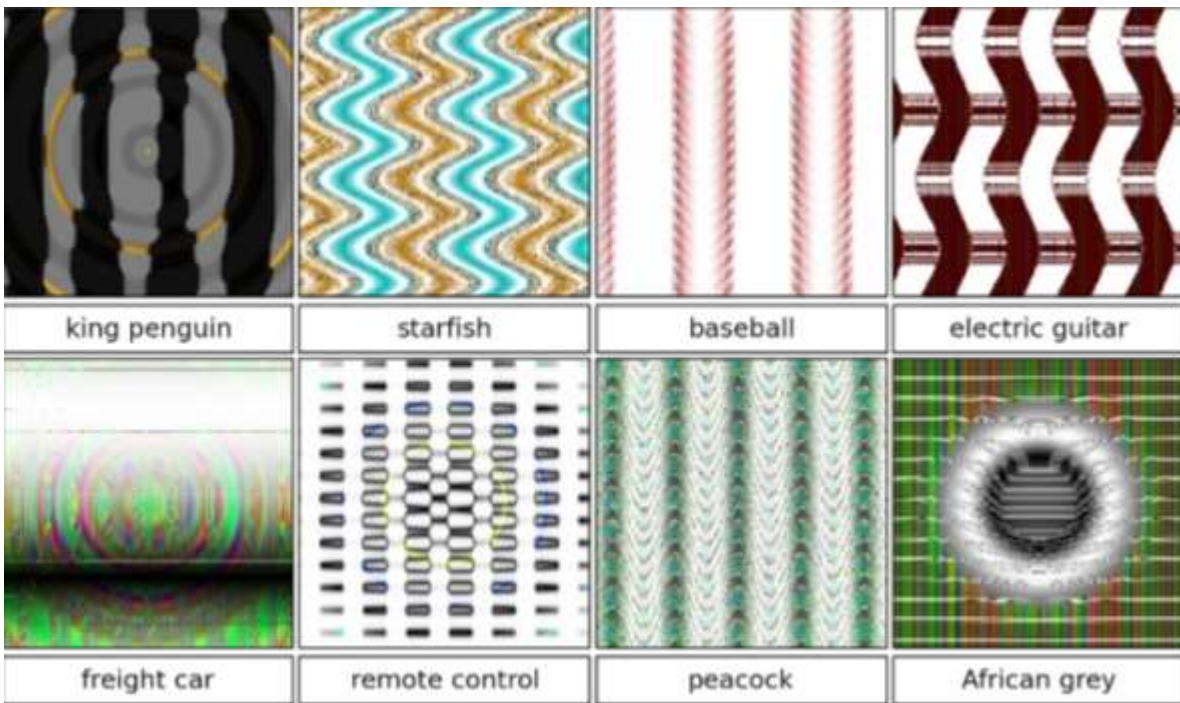
99.99%  
ostrich, *Struthio camelus*



# Why Gen. Models?

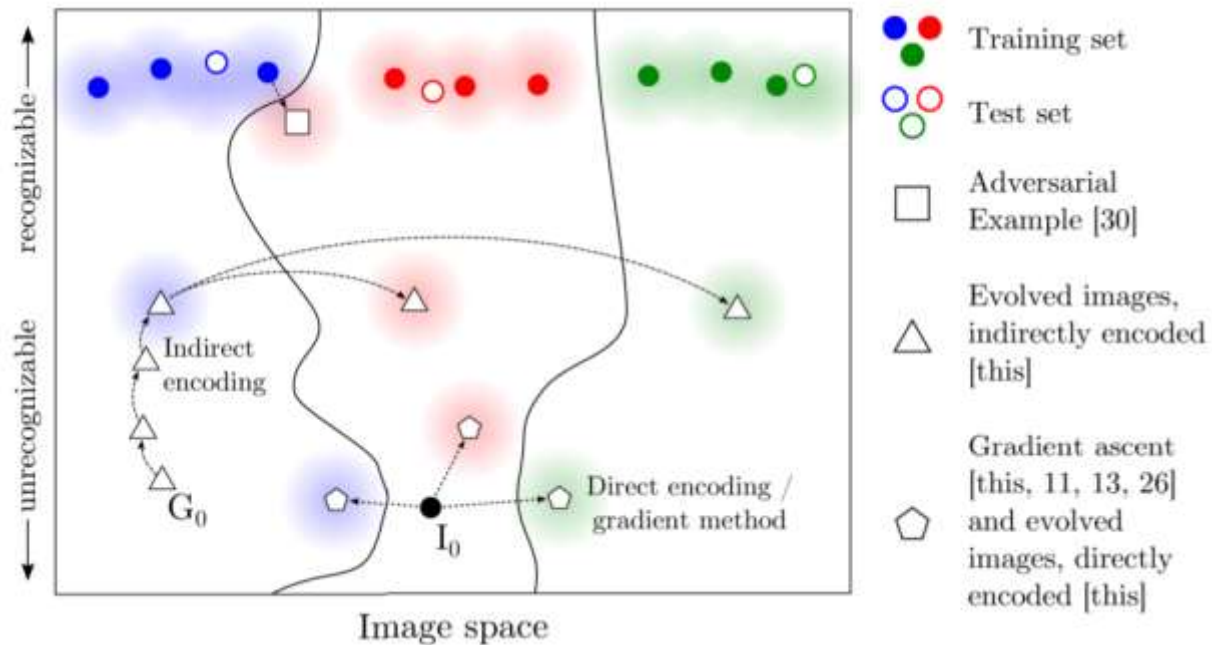


# Why Gen. Models?





# Why Gen. Models?

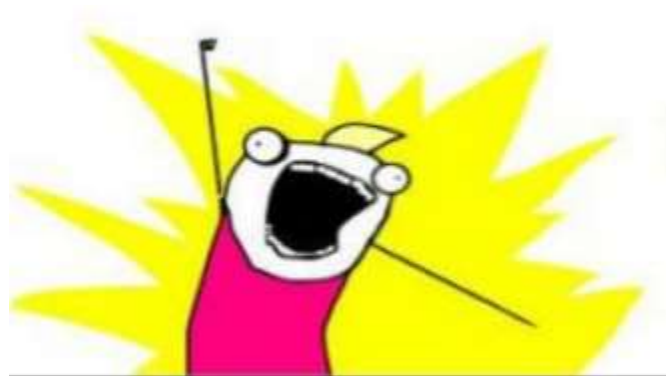


# What do we want?

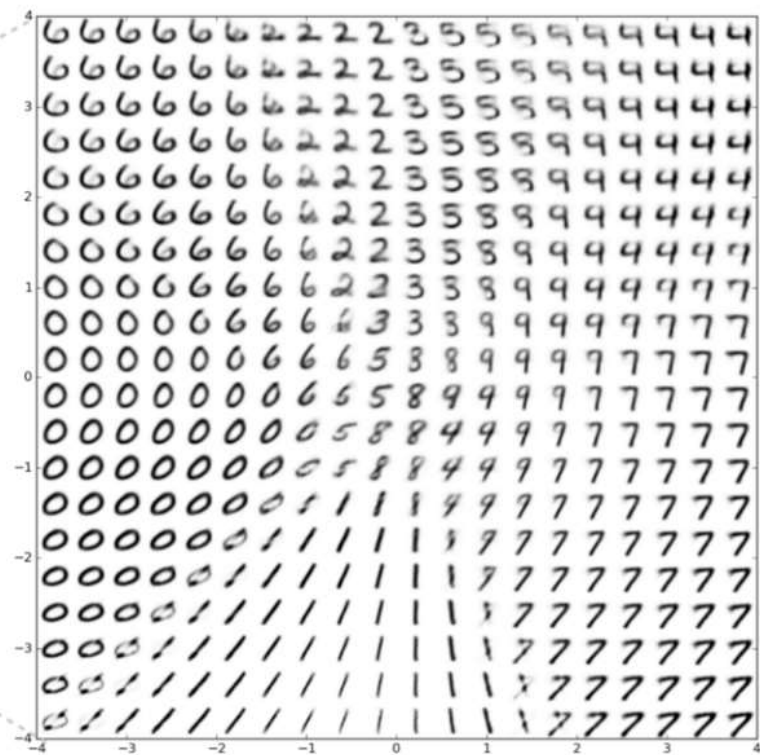
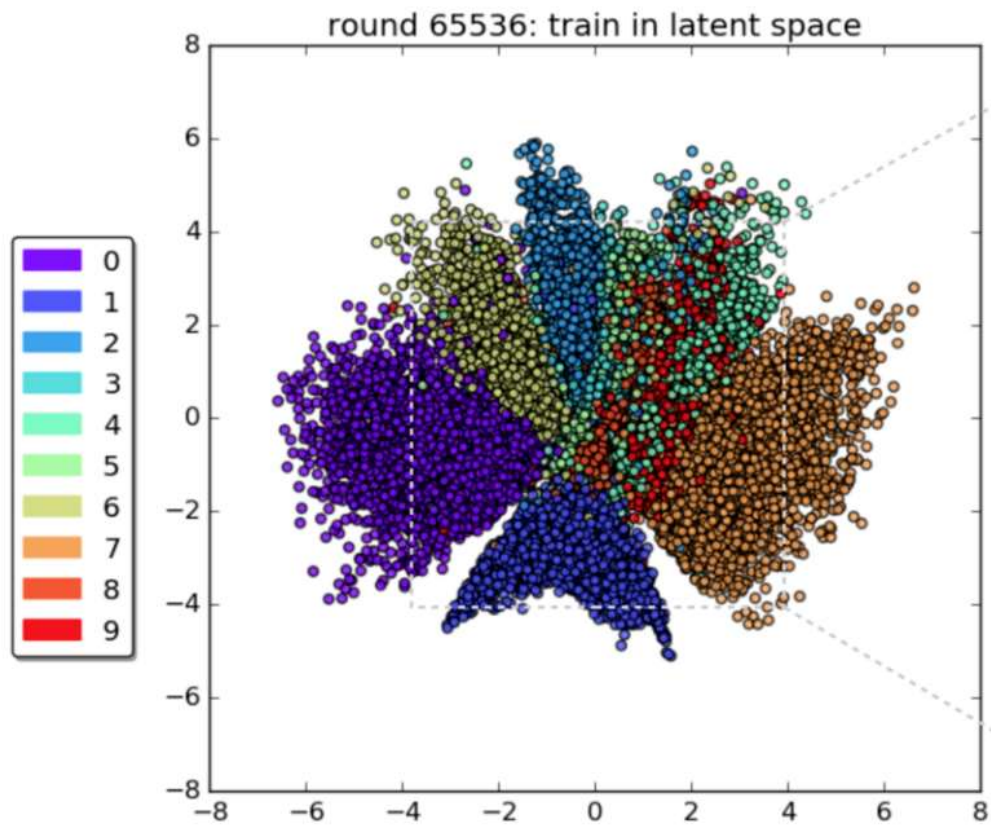
- Generative model
- “Structure constraint” on latent space

## Why?

- Semi-Supervised learning
- Visualize z-space
- Not so easily fooled
- More...



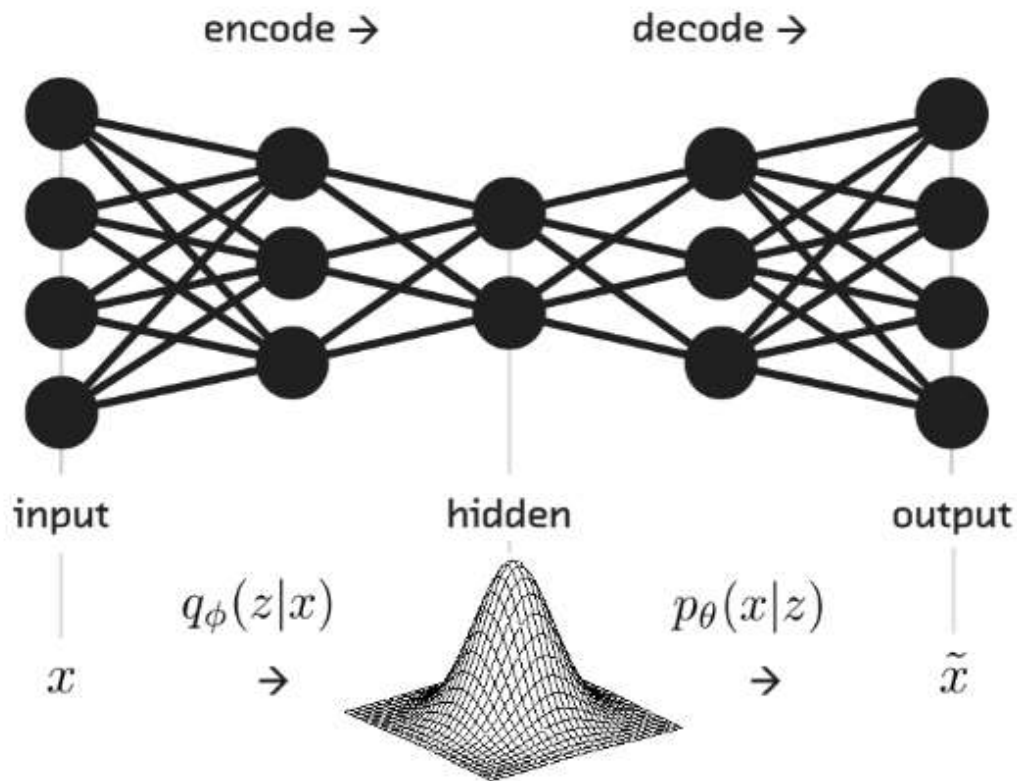




[illegible]

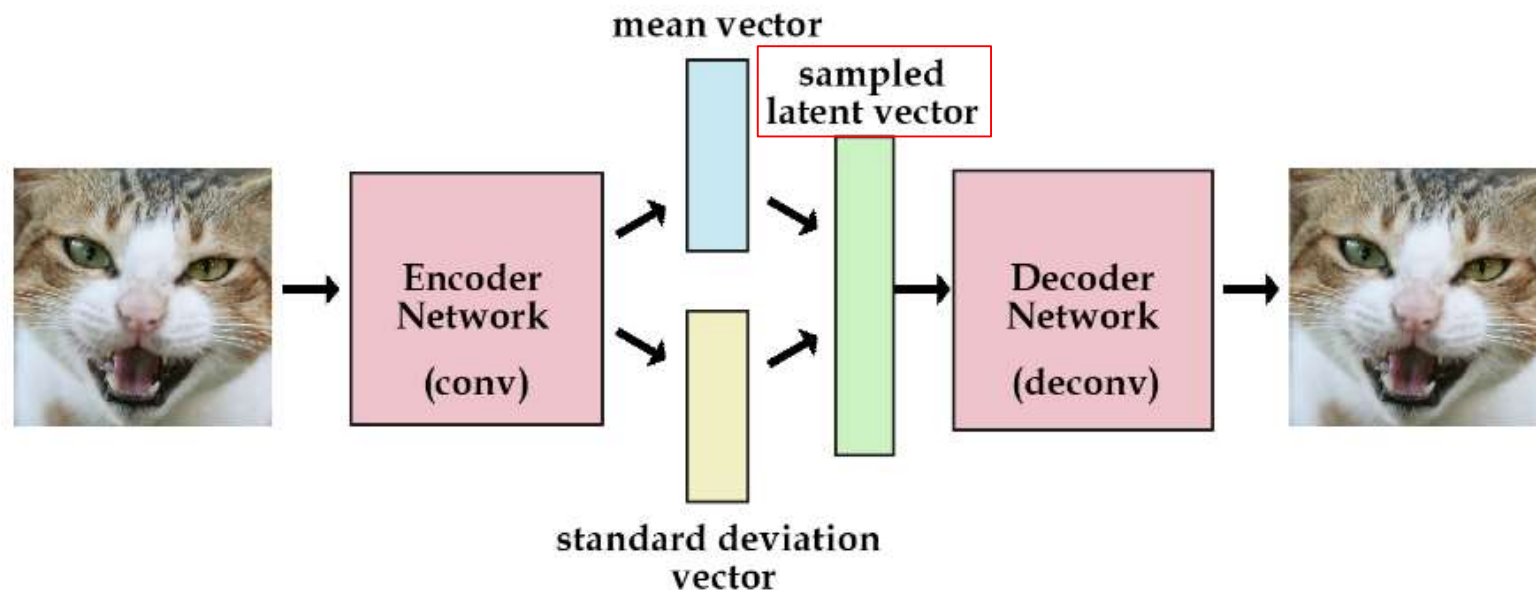
# AutoEncoder Attempt #1

- Encoder  
 $q(z|x)$ : get  $z$  given  $x$
- Decoder  
 $p(x|z)$ : get  $x$  given  $z$
- What's the difference?



# AutoEncoder Attempt #1

Problems?

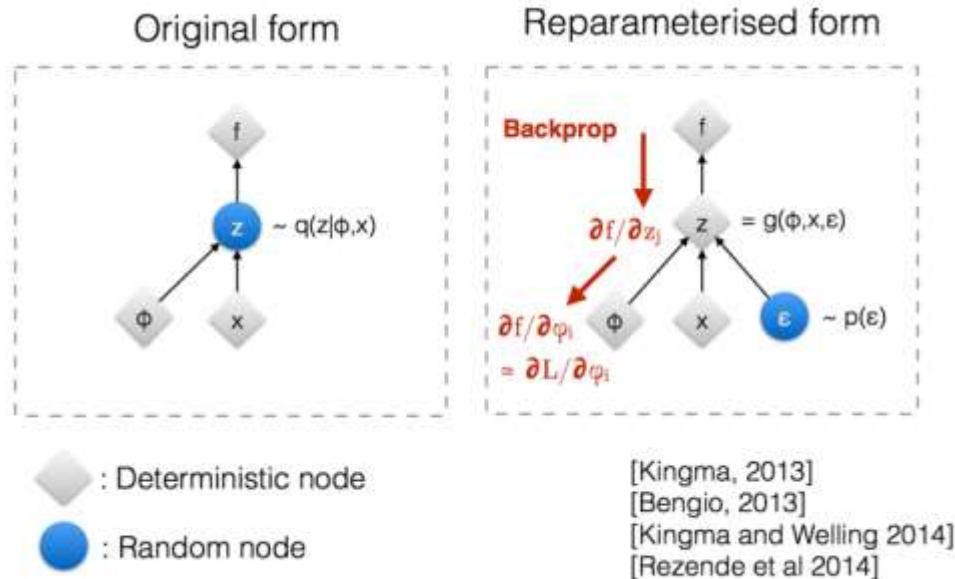


# Re-Parameterization Trick

Backpropagation not possible through random sampling!

$$z^{(i,l)} = \mu^{(i)} + \sigma^{(i)} \odot \epsilon_i$$

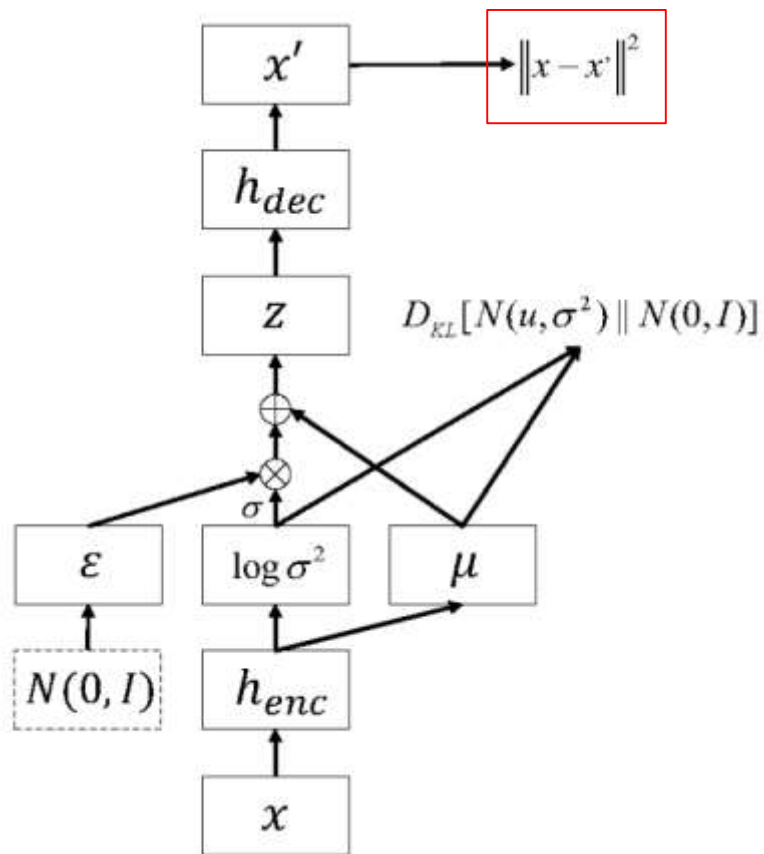
$$\epsilon_i \sim N(0,1)$$



[<https://arxiv.org/abs/1609.04468>]

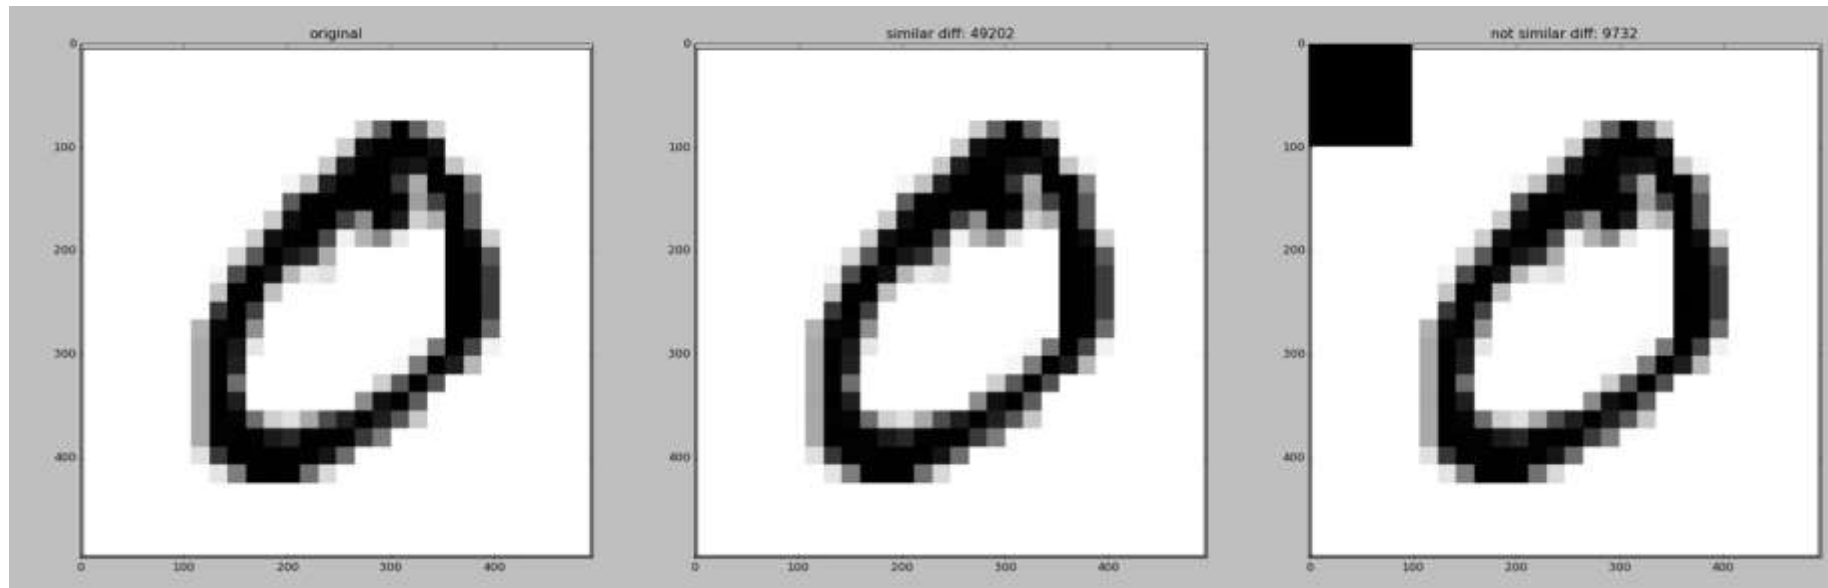
# Net Structure

- More problems?





# L2 distance in pixel space



# GAN - Generative Adversarial Networks

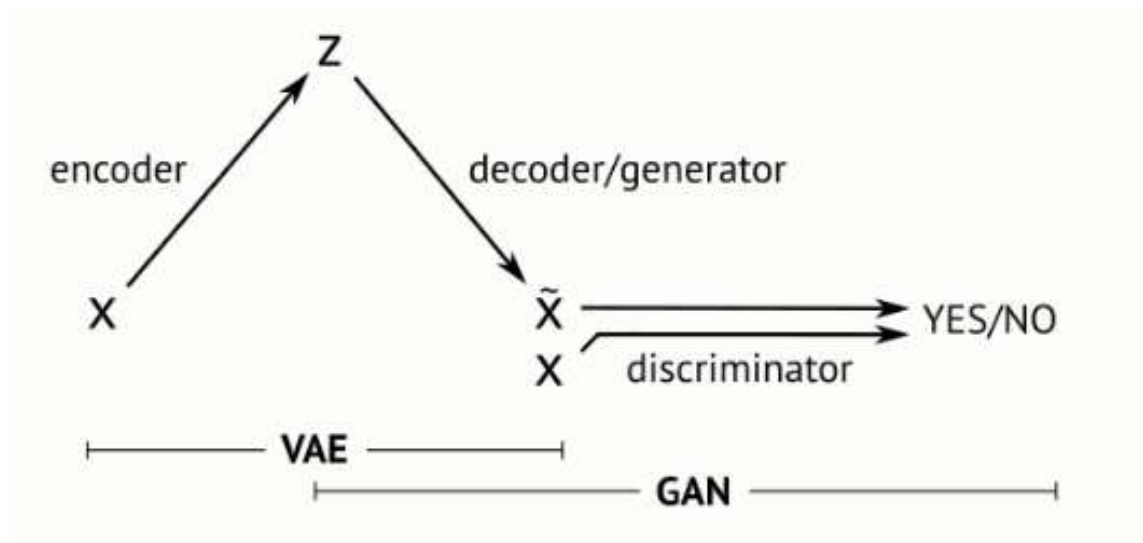
“You pit a generative (G) machine against a discriminative (D) machine and make them fight.”

© Soumith Chintala

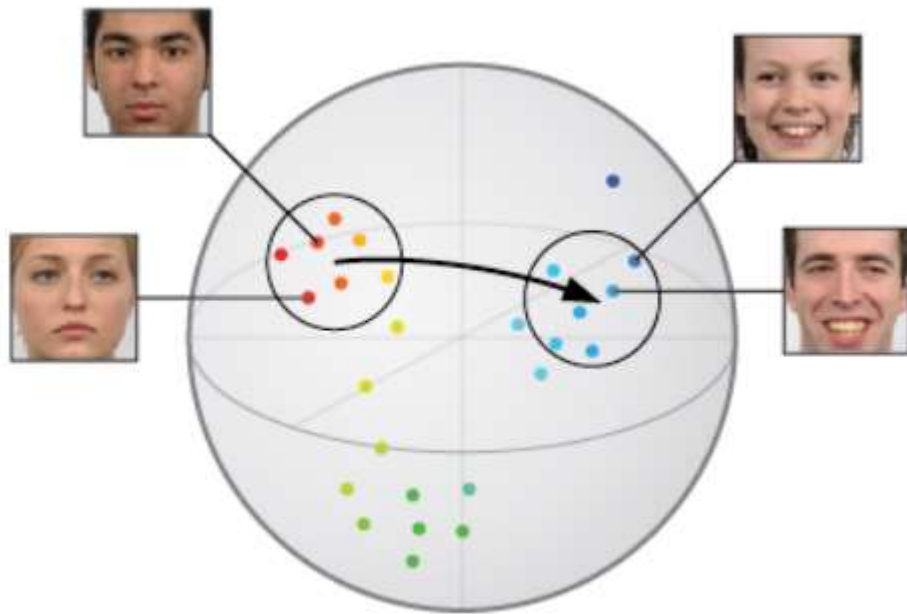
<http://soumith.ch/eyescream/>



# GAN - learn the loss

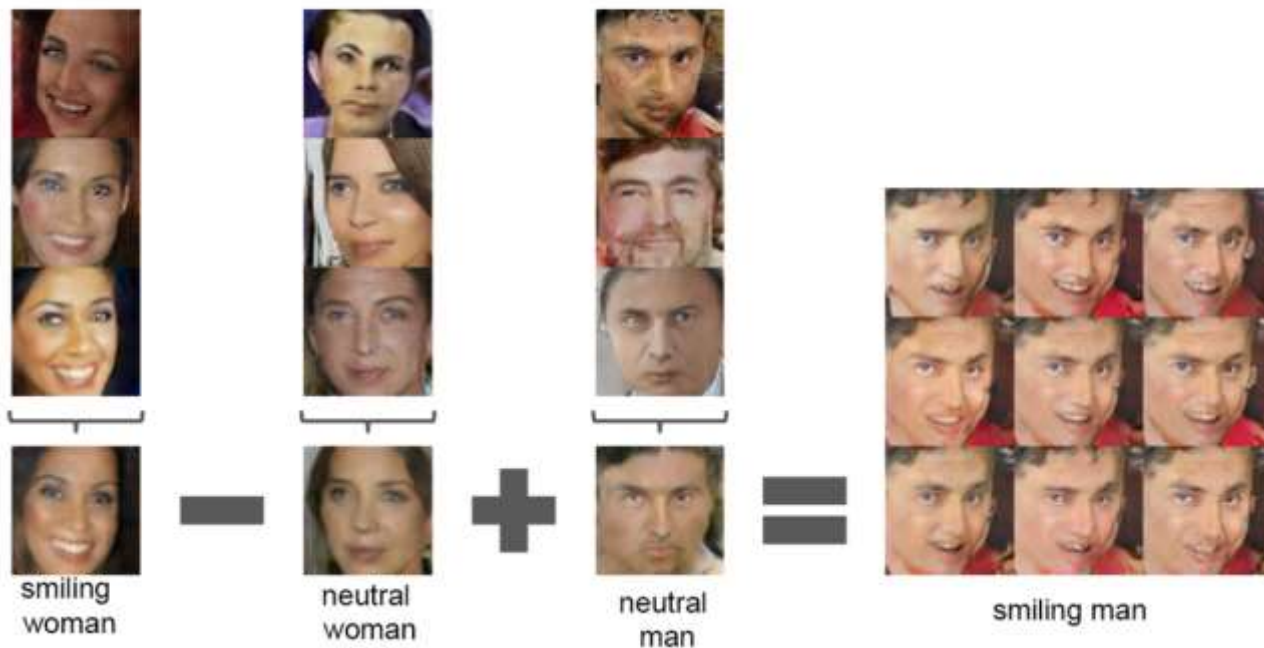


If you do it right!



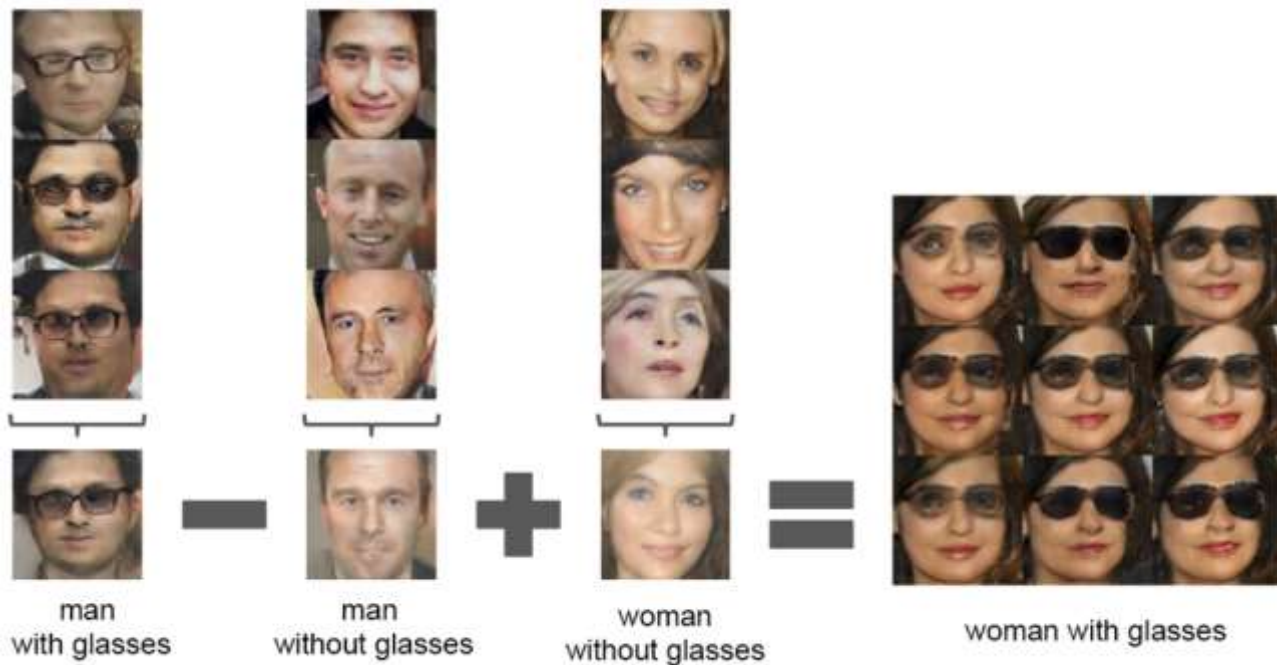
# If you do it right!

Arithmetic on faces



# If you do it right!

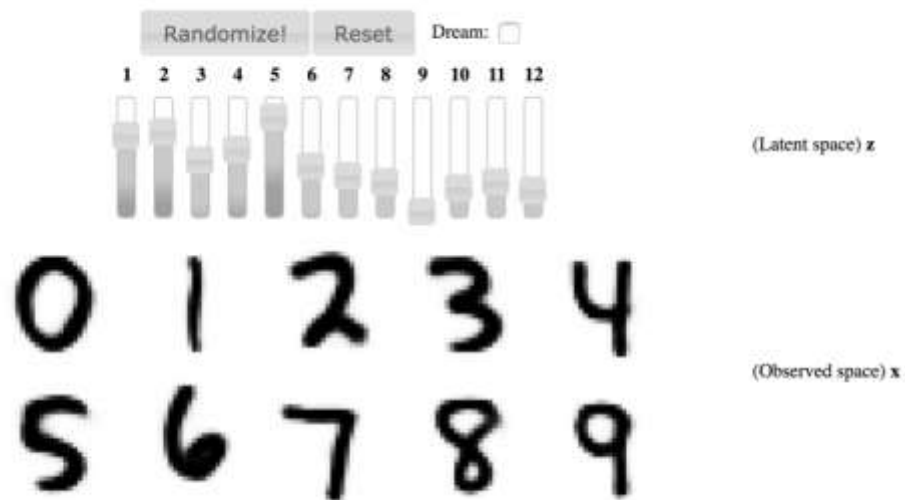
↻ Arithmetic on faces





# If you do it right!

[http://www.dpkgingma.com/sgvb\\_mnist\\_demo/demo.html](http://www.dpkgingma.com/sgvb_mnist_demo/demo.html)

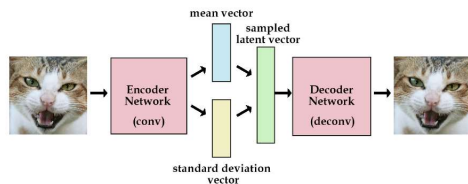


# Take Aways

- Employ Structure, It's cool
- GAN may be your next loss function
  - Super Res
  - Pixel Level Seg.
  - AutoEncoder (we saw it today)
- Re-Parameterization Trick

----Personal takeaw

- Don't give up when



1000)

# References

- [https://github.com/oduerr/dl\\_tutorial/blob/master/tensorflow/vae/vae\\_demo-2D.ipynb](https://github.com/oduerr/dl_tutorial/blob/master/tensorflow/vae/vae_demo-2D.ipynb)
- <https://home.zhaw.ch/~dueo/bbs/files/vae.pdf>
- <http://kvfrans.com/variational-autoencoders-explained/>
- <http://blog.fastforwardlabs.com/2016/08/12/introducing-variational-autoencoders-in-prose-and.html>
- <http://blog.fastforwardlabs.com/2016/08/22/under-the-hood-of-the-variational-autoencoder-in.html>
- [https://github.com/Newmu/dcgan\\_code](https://github.com/Newmu/dcgan_code)
- <http://torch.ch/blog/2015/11/13/gan.html>
- <https://github.com/soumith/ganhacks>
- <https://research.fb.com/wp-content/uploads/2016/11/luc16wat.pdf>
- <https://arxiv.org/pdf/1412.1897.pdf>



# Re-Parameterization Trick

Sampling (reparametrization trick)

$$z^{(i,l)} \sim N(\mu^{(i)}, \sigma^{2(i)})$$

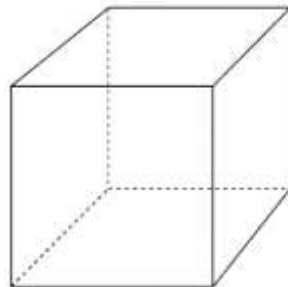
$$z^{(i,l)} = \mu^{(i)} + \sigma^{(i)} \odot \varepsilon_i \quad \varepsilon_i \sim N(0,1)$$

Writing  $z$  in this form, results in a deterministic part and noise.

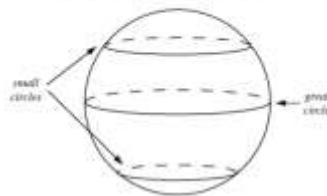
Cannot back propagate through a random drawn number

$z$  has the same distribution, but now one can back propagate.

- Don't sample from a Uniform distribution



- Sample from a gaussian distribution



[<https://arxiv.org/abs/1609.04468>]